


Collaboration & Competition

1. Introduction

The goal of this project is to train 2 agents in the Unity Tennis environment to play tennis. In this environment, two agents control rackets to bounce a ball over a net. If an agent hits the ball over the net, it receives a reward of $+0.1$. If an agent lets a ball hit the ground or hits the ball out of bounds, it receives a reward of -0.01 . Thus, the goal of each agent is to keep the ball in play.




The observation space consists of 8 variables corresponding to the position and velocity of the ball and racket. Each agent receives its own, local observation. Two continuous actions are available, corresponding to movement toward (or away from) the net, and jumping.

The task is episodic, and in order to solve the environment, the agents must get an average score of $+0.5$ over 100 consecutive episodes, after taking the maximum over both agents.

2. Algorithm

The agent is trained with the Multi Agent DDPG. The full algorithm is described in the methods section of the paper¹. The information flow in a MADDPG algorithm is described below.



The environment was solved by tuning the DDPG algorithm described in the second project of this course (see continuous control). In detail, an agent called MADDPG (Multi-Agent-DDPG) handles the training of two DDPG agents, each one of them corresponding to one of the two players. This MADDPG agent handles a common replay buffer for its two sub-agents.

Each one the DDPG agent has:

- An actor seeing the state corresponding to its player and performing an action for it
- A critic seeing the whole picture: The state of both agents, and estimating the action-value function.

The training procedure goes like this:

- States, actions, rewards, next states are collected in a single replay buffer.
- Each agent samples a batch of experiences from the replay buffer and asks the other agent to propose an action for each state and next states
- Each agent is trained like DDP agents, with minor differences. The critic takes as input the states of all agents. The best actions and best next actions are estimated thanks to both agents.

¹ <https://arxiv.org/abs/1706.02275>

3. Hyperparameters and Architecture

Hyperparameters

```

BUFFER_SIZE = int(1e6) # replay buffer size
BATCH_SIZE = 128      # minibatch size
LR_ACTOR = 1e-3        # learning rate of the actor
LR_CRITIC = 1e-3       # learning rate of the critic
WEIGHT_DECAY = 0       # L2 weight decay
LEARN_EVERY = 1        # learning timestep interval
LEARN_NUM = 1          # number of learning passes
GAMMA = 0.99           # discount factor

TAU = 7e-2             # for soft update of target parameters
OU_SIGMA = 0.2         # Ornstein-Uhlenbeck noise parameter, volatility
OU_THETA = 0.12        # Ornstein-Uhlenbeck noise parameter, speed of
                        # mean reversion
EPS_START = 5.5        # initial value for epsilon in noise decay process in
                        # Agent.act()
EPS_EP_END = 250       # episode to end the noise decay process
EPS_FINAL = 0          # final value for epsilon after decay

```

Actor neural network

The Actor neural network consists of two fully connected layers.

- The input has 24 units to observe a state with a dimension of 24.
- The first hidden layer has a size of 256 units and the second layer layer has 128 units.
- The output has 2 units for one dimension movement and jumping.

Critic neural network

The Critic neural network consists of two fully connected layers.


- The input has 24 units to observe a state with a dimension of 24.
- The first hidden layer has a size of 256 units and the second layer layer has 128 units + 2 (actions).
- The output has 1 unit.

3. Results

Given the chosen architecture and parameters, our results are:

Episode 10	Max: 0.400000	Average: 0.190000
Episode 100	Max: 1.000000	Average: 0.311000
Episode 200	Max: 0.500000	Average: 0.254000
Episode 300	Max: 0.500000	Average: 0.252900
Episode 400	Max: 0.800000	Average: 0.322000
Episode 500	Max: 0.400000	Average: 0.327000
Episode 600	Max: 0.400000	Average: 0.305000
Episode 700	Max: 0.600000	Average: 0.303000
Episode 800	Max: 0.600000	Average: 0.308900
Episode 900	Max: 0.500000	Average: 0.429000
Episode 1000	Max: 0.700000	Average: 0.382000
Episode 1100	Max: 0.500000	Average: 0.364000
Episode 1200	Max: 0.700000	Average: 0.321000
Episode 1300	Max: 0.600000	Average: 0.382000
Episode 1400	Max: 0.800000	Average: 0.388000
Episode 1410	Max: 0.500000	Average: 0.388900
Episode 1420	Max: 1.700000	Average: 0.414900
Episode 1430	Max: 2.200000	Average: 0.455900
Episode 1440	Max: 0.700000	Average: 0.457900
Episode 1450	Max: 0.700000	Average: 0.471800
Episode 1460	Max: 0.600000	Average: 0.486800
Episode 1470	Max: 0.600000	Average: 0.483800

Environment solved in 1375 episodes. Average: 0.503800 over past 100 episodes



These results meets the project's requirements as the agents are able to receive an average score of 0.504 over the last 100 episodes, from episode 1375.

4. Ideas for Future Work

Further improvements could be:

- Further tune the hyper-parameters to converge faster, automated grid based
- Implement a prioritized replayed-buffer
- Train two pairs of agents, and then switch the partners, to see if they are able to play with another player.